

# The Virtual World




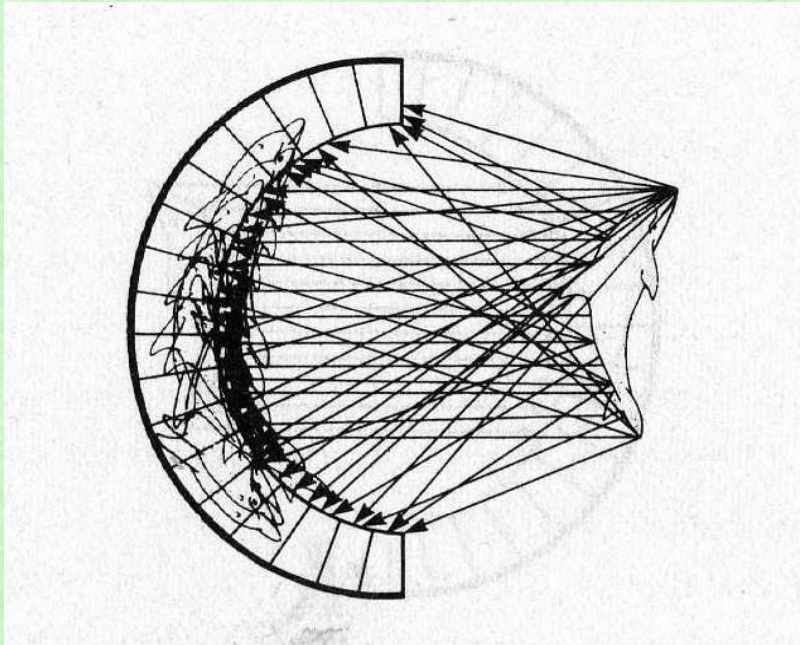
# Building a Virtual World

- Goal: mimic human vision in a virtual world (with the computer)
  - Cheating for efficiency, using knowledge about light and our eye (e.g. from the last lecture)
- Create a virtual **camera**, place it somewhere, and point it at something
- Put **film** containing **pixels** with **RGB values** ranging from **0-255** into the camera
  - Taking a picture creates film data as the final image
- Place **objects** into the world, including a floor/ground, walls, ceiling/sky, etc.
  - Two step process: (1) make objects (geometric modeling), (2) place objects (transformations)
  - Making objects is itself a two step process: (1) build geometry (geometric modeling), (2) paint geometry (texture mapping)
- Put **lights** into the scene (so that it's not completely dark)
- Finally, snap the picture:
  - “Code” emits light from virtual light sources, bounces that light off of virtual geometry, and follows that light into the camera and onto the film
  - We will consider both scanline rendering and ray tracing for the taking this picture



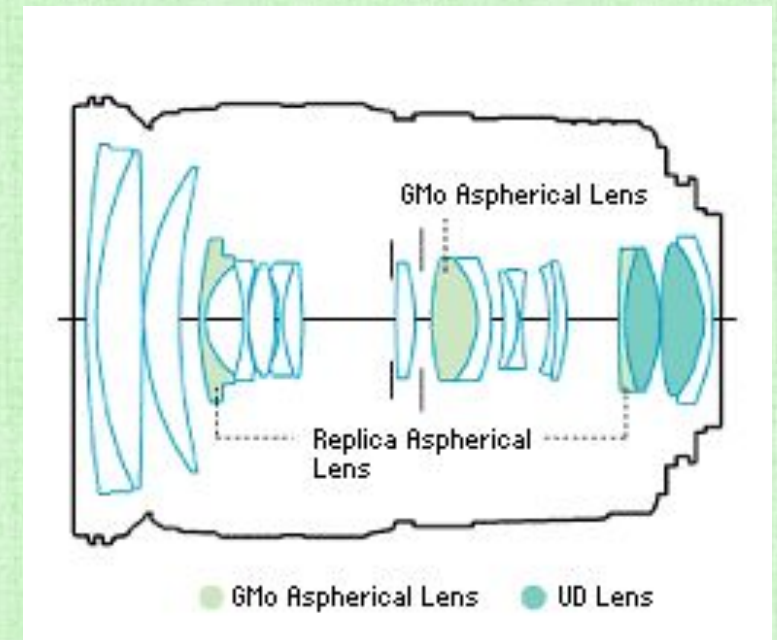
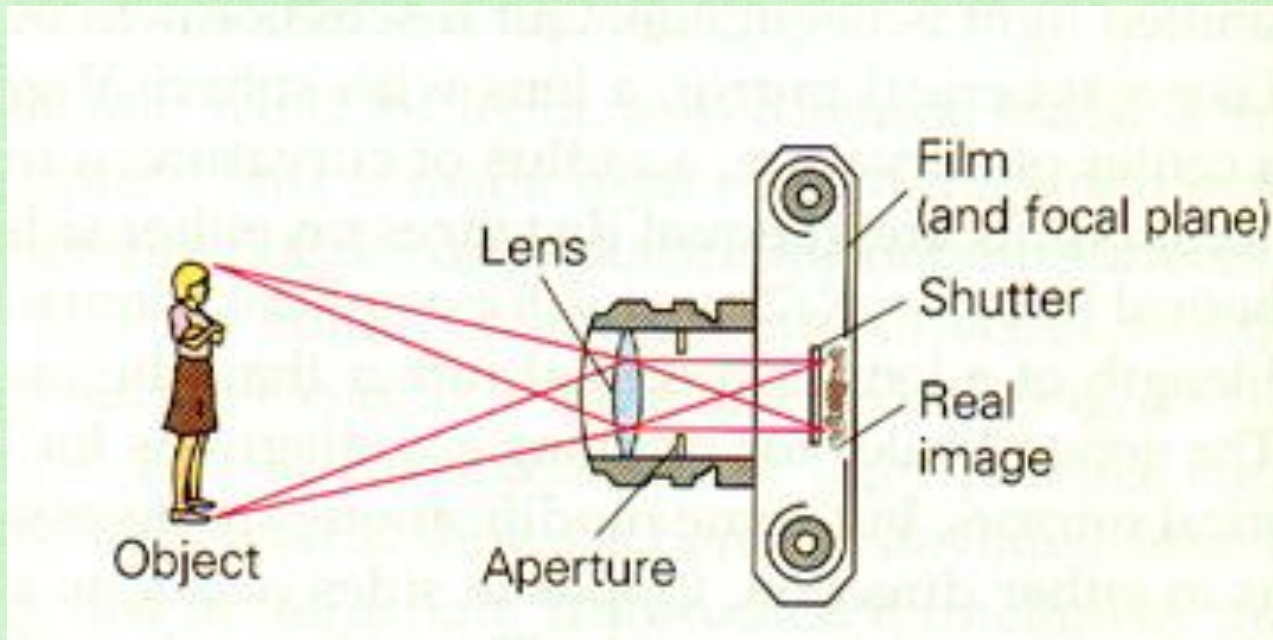
# Pupil

- Light emanates off of every point of an object outwards in every direction
  - That's why we can all look at the same spot on the same object and see it
  - Light is leaving from that point/spot on the object and entering each of our eyes
- Without a pupil, light from every part of an object would hit the same cone on our eye, blurring the image
- The (small) pupil restricts the entry of light so that each cone only receives light emanating from a small region on the object, giving interpretable spatial detail



# Aperture

- Cameras are similar to the eye, except with mechanical as opposed to biological components
- Instead of cones, the camera has mechanical pixels
- Instead of a pupil, the camera has a small (adjustable) aperture that light passes through
- The camera also typically has a hefty/complex lens system





# Aside: Lens Flare

- Many camera complexities are not often properly accounted for in virtual worlds
- Thus, some complex effects like depth of field, motion blur, chromatic aberration, lens flare, etc. have to be approximated/modeled in other ways (as we will discuss later)
- Particularly complex is lens flare, which is caused by light being reflected/scattered by lenses in the lens system
- This is caused in part by material inhomogeneities in the lens, and depends on the geometry of lens surfaces and characteristic planes, absorption/dispersion of lens elements, antireflective coatings, diffraction, etc.



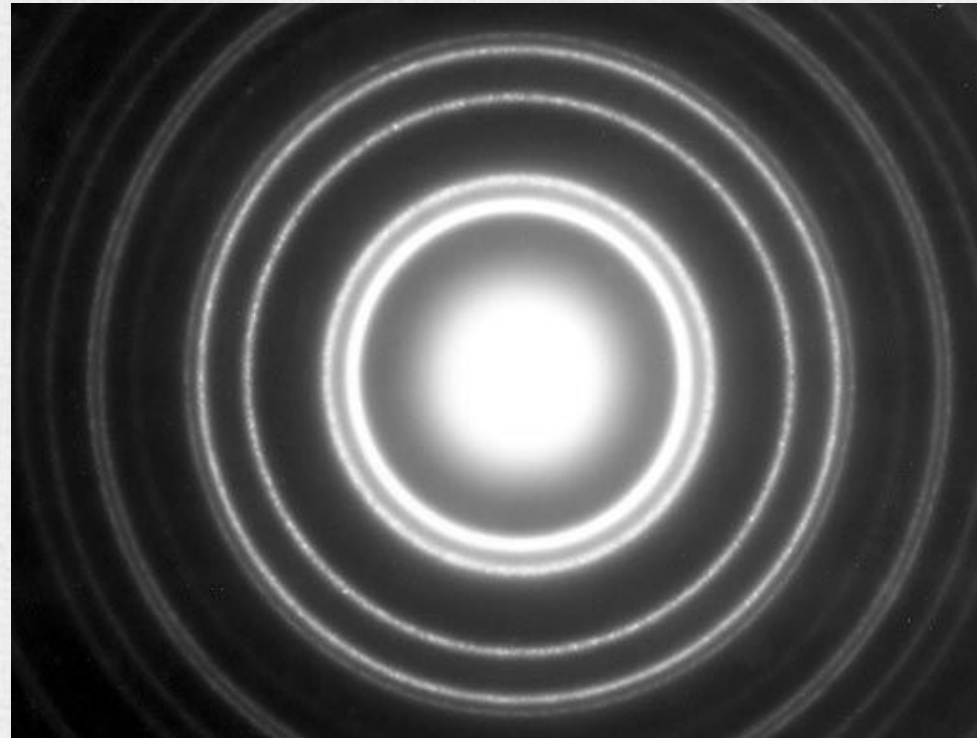
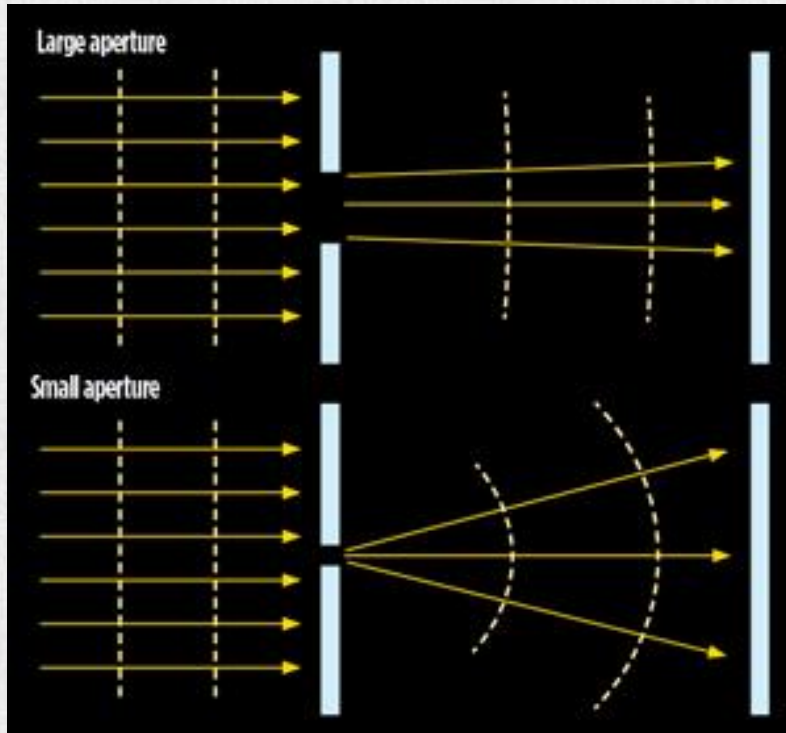
# Pinhole Camera

- The pupil/aperture has to have a finite size in order for light to get through
- If too small, not enough light enters and the image is too dark/noisy to interpret
  - In addition, light can diffract (instead of traveling in straight lines) distorting the image
- If too large, light from a large area of an object hits the same cone causing blurring
- Luckily, the virtual world can use a single point for the aperture (without worrying about dark or distorted images)



# Aside: Diffraction


- Light spread out as it goes through small openings
- This happens when the camera aperture is too small (diffraction limited)
- Creates constructive/destructive interference of light waves (the Airy disk effect)





# Pinhole Camera

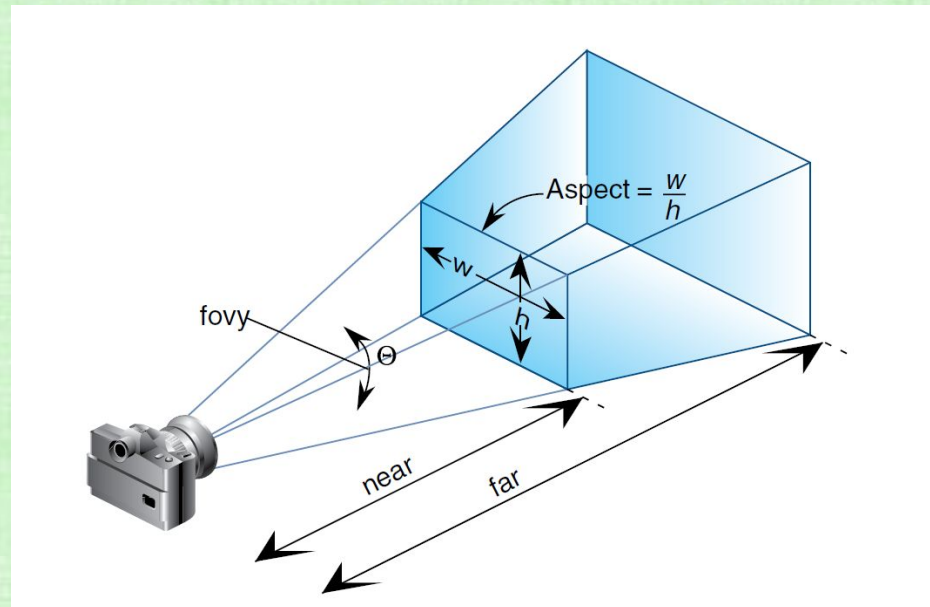
- Light leaving any point travels in straight lines
- We only care about the lines that hit the pinhole (a single point)
- Infinite depth of field; i.e., everything is in focus (no blurring)
- An upside down image is formed by the intersection of these lines with an image plane
- More distant objects subtend smaller visual angles and appear smaller
- Objects occlude objects behind them





# Virtual World Camera

- Trick: Move the film out in front of the pinhole, so that the image is not upside down
- Only render (compute an image for) objects further away from the camera than the film
- Add a back clipping plane for efficiency
- Volume between the film (front clipping plane) and the back clipping plane is the viewing frustum (shown in blue)
  - Make sure near/far (i.e. front/back) clipping planes have enough space between them to contain the scene
  - Make sure objects are inside the viewing frustum
  - Do not set the near clipping plane at the camera aperture!



# Cameras Distortion depends on Distance

- Do not put the camera too close to objects of interest!
  - Significant/severe deductions for poor camera placement, fisheye, etc. (because the image will look terrible)
- Set up the scene like a real world scene!
- Get very familiar with the virtual camera!




@160CM



@25CM

# Eye Distortion?

- Your eye also has distortion
  - Unlike a camera, you don't actually see the signal received on the cones
  - Instead, you perceive an image (highly) processed by your brain
  - Your eyes constantly move around obtaining multiple images for your brain to work with
  - You have two eyes, and see in stereo, so triangulation can be used to estimate depth and undo distortion
- 
- If your skeptical about all this processing, remember that your eye sees this:



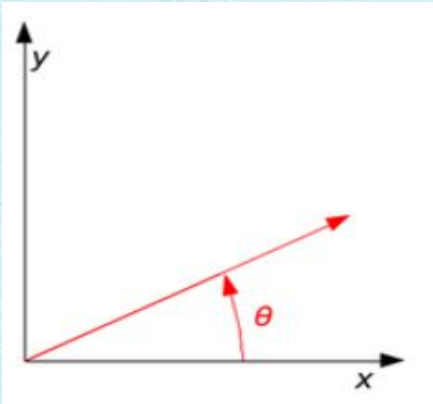


# Dealing with Objects

- 
- Let's start with a single 3D point  $\vec{x} = \begin{pmatrix} x \\ y \\ z \end{pmatrix}$  and move it around in the virtual world
- An object is just a collection of points, and as such the methods for handling a single point readily extend to handling entire objects
- Typically, objects are created in a reference space, which we refer to as object space
- After creation, we place objects into the scene, which we refer to as world space
- This may requires **rotation**, **translation**, **resizing** of the object
- When taking a picture, points on the object are projected onto the film, which we refer to as screen space
- Unlike rotation/translation/resizing, the projection onto screen space is highly nonlinear and the source of undesirable distortion

# Rotation

- 
- Consider a single 3D point,  $\vec{x} = \begin{pmatrix} x \\ y \\ z \end{pmatrix}$
- In 2D, one can rotate a point clockwise about the origin via:



$$\begin{pmatrix} x^{new} \\ y^{new} \end{pmatrix} = \begin{pmatrix} \cos \theta & -\sin \theta \\ \sin \theta & \cos \theta \end{pmatrix} \begin{pmatrix} x \\ y \end{pmatrix} = R(\theta) \begin{pmatrix} x \\ y \end{pmatrix}$$

$$R_z(\theta) = \begin{pmatrix} \cos \theta & -\sin \theta & 0 \\ \sin \theta & \cos \theta & 0 \\ 0 & 0 & 1 \end{pmatrix}$$



# Rotation

- To rotate a 3D point around the x-axis, y-axis, or z-axis (respectively), one multiplies by:

$$R_x(\theta) = \begin{pmatrix} 1 & 0 & 0 \\ 0 & \cos \theta & -\sin \theta \\ 0 & \sin \theta & \cos \theta \end{pmatrix}$$

$$R_y(\theta) = \begin{pmatrix} \cos \theta & 0 & \sin \theta \\ 0 & 1 & 0 \\ -\sin \theta & 0 & \cos \theta \end{pmatrix}$$

$$R_z(\theta) = \begin{pmatrix} \cos \theta & -\sin \theta & 0 \\ \sin \theta & \cos \theta & 0 \\ 0 & 0 & 1 \end{pmatrix}$$

- Matrix multiplication doesn't commute, i.e.  $AB \neq BA$ , so the **order** of rotations **matters**!
- Rotating about the x-axis and then the y-axis,  $R_y(\theta_y)R_x(\theta_x)\vec{x}$ , gives a different results than rotating about the y-axis and then the x-axis,  $R_x(\theta_x)R_y(\theta_y)\vec{x}$



# Line Segments are Preserved

- Consider two points  $\vec{p}$  and  $\vec{q}$  and the line segment between them:

$$\vec{u}(\alpha) = (1 - \alpha)\vec{p} + \alpha\vec{q}$$

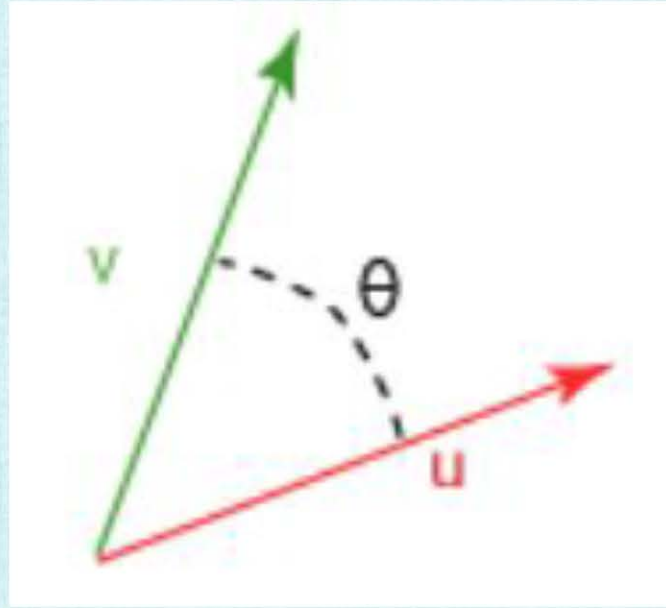
- Here,  $\vec{u}(0) = \vec{p}$  and  $\vec{u}(1) = \vec{q}$  and  $0 \leq \alpha \leq 1$  determine all the points on the line segment
- Multiplying points on the line segment by a rotation matrix  $R$  gives:

$$R\vec{u}(\alpha) = R((1 - \alpha)\vec{p} + \alpha\vec{q}) = (1 - \alpha)R\vec{p} + \alpha R\vec{q}$$

- Here,  $R\vec{u}(0) = R\vec{p}$  and  $R\vec{u}(1) = R\vec{q}$  and  $0 \leq \alpha \leq 1$  determine all the points on the new rotated line segment connecting  $R\vec{p}$  and  $R\vec{q}$ 
  - i.e., one only need rotate the endpoints to construct the new line segment connecting them
- $\|R\vec{u}(\alpha) - R\vec{p}\|_2^2 = \|R(\vec{u}(\alpha) - \vec{p})\|_2^2 = (\vec{u}(\alpha) - \vec{p})^T \mathbf{R}^T \mathbf{R} (\vec{u}(\alpha) - \vec{p}) = \|\vec{u}(\alpha) - \vec{p}\|_2^2$  shows that the distance between rotated points is equivalent to the distance between original points

# Angles are Preserved

- Consider two line segments  $\vec{u}$  and  $\vec{v}$  with  $\vec{u} \cdot \vec{v} = \|\vec{u}\|_2 \|\vec{v}\|_2 \cos(\theta)$  where  $\theta$  is the angle between them




- $R\vec{u} \cdot R\vec{v} = \vec{u}^T \mathbf{R}^T \mathbf{R} \vec{v} = \vec{u}^T \vec{v} = \|\vec{u}\|_2 \|\vec{v}\|_2 \cos(\theta) = \|R\vec{u}\|_2 \|R\vec{v}\|_2 \cos(\theta)$
- So, the angle  $\theta$  between  $\vec{u}$  and  $\vec{v}$  is the same as the angle  $\theta$  between  $R\vec{u}$  and  $R\vec{v}$



# Shape is Preserved

- In continuum mechanics, one measures the deformation of a material by a tensor called the strain
- The six unique entries in the nonlinear Green strain tensor are computed by comparing an undeformed tetrahedron to its deformed counterpart
- Given a tetrahedron in 3D, it is fully determined by one point and three line segments (the dotted lines in the figure)

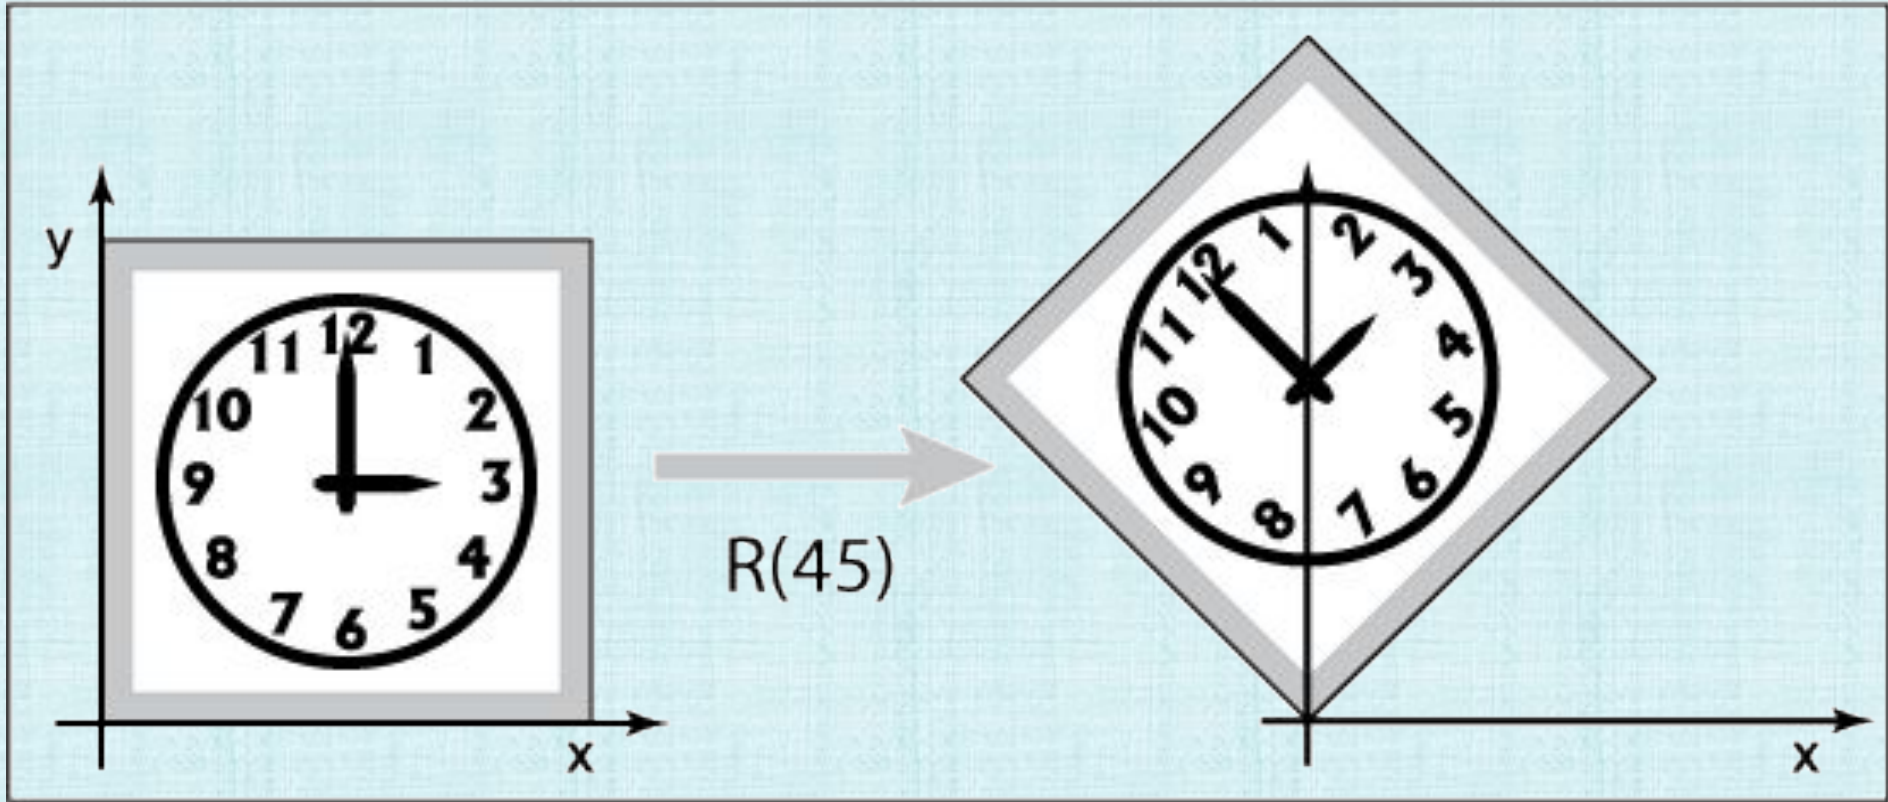


- The 3 lengths of these three line segments and the 3 angles between any two of them are used to compare the undeformed tetrahedron to its deformed counterpart
- Since we proved these were all identical under rotations, rotations are shape preserving



# Shape is Preserved

- Thus we can rotate entire objects without changing them

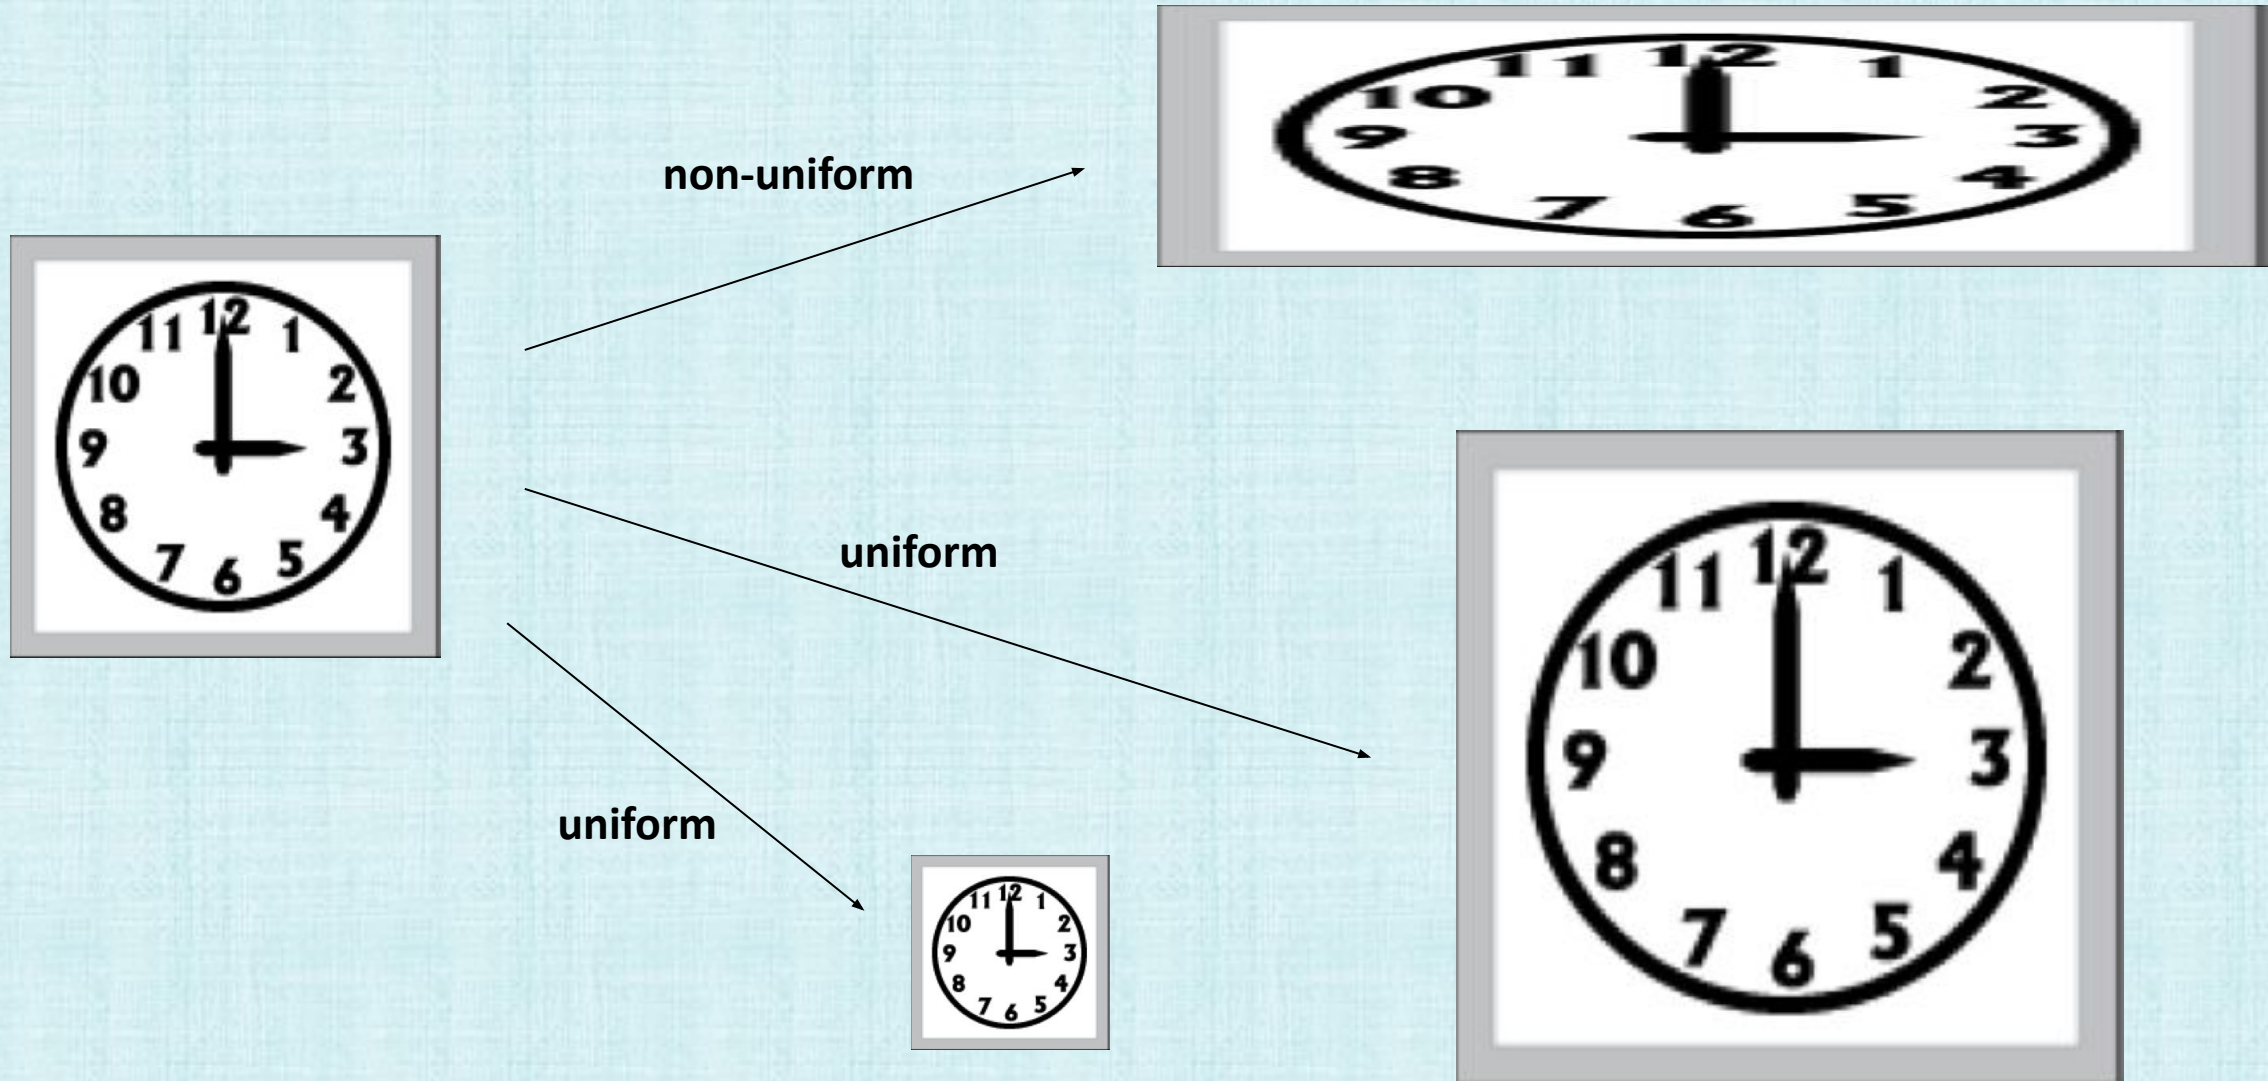


# Scaling (or Resizing)

- A scaling matrix has the form  $S = \begin{pmatrix} s_1 & 0 & 0 \\ 0 & s_2 & 0 \\ 0 & 0 & s_3 \end{pmatrix}$  and can both scale and shear the object
- Generally speaking, shearing an object changes lengths/angles creating significant distortion
- When  $s_1 = s_2 = s_3$ , one has pure scaling of the form  $S = \begin{pmatrix} s & 0 & 0 \\ 0 & s & 0 \\ 0 & 0 & s \end{pmatrix} = sI$
- The distributive law of matrix multiplication guarantees that line segments map to line segments, and  $\|S\vec{u}(\alpha) - S\vec{p}\|_2^2 = s\|\vec{u}(\alpha) - \vec{p}\|_2^2$  implies that the distance between scaled points is increased/decreased by a factor of  $s$
- $S\vec{u} \cdot S\vec{v} = s^2\vec{u} \cdot \vec{v} = s^2\|\vec{u}\|_2 \|\vec{v}\|_2 \cos(\theta) = \|S\vec{u}\|_2 \|S\vec{v}\|_2 \cos(\theta)$  shows that angles between line segments are preserved
- Thus, when using uniform scaling, objects grow/shrink but look the same as far as proportions are concerned (they are mathematically similar)



# Scaling (or Resizing)





# Homogenous Coordinates

- In order to deal with transformations via matrix multiplication, one uses homogeneous coordinates
- The homogeneous coordinates of a 3D point  $\vec{x} = \begin{pmatrix} x \\ y \\ z \end{pmatrix}$  are  $\vec{x}_H = \begin{pmatrix} xw \\ yw \\ zw \\ w \end{pmatrix}$  for any  $w \neq 0$
- Dividing any homogenous coordinates by its fourth component gives  $\begin{pmatrix} x \\ y \\ z \\ 1 \end{pmatrix}$  or  $\begin{pmatrix} \vec{x} \\ 1 \end{pmatrix}$
- We convert all our 3D points to the form  $\vec{x}_H = \begin{pmatrix} x \\ y \\ z \\ 1 \end{pmatrix}$ , i.e.  $w = 1$ , to deal with translations
- For vectors  $\vec{u} = \begin{pmatrix} u_1 \\ u_2 \\ u_3 \end{pmatrix}$ , the homogenous coordinates are  $\vec{u}_H = \begin{pmatrix} u_1 \\ u_2 \\ u_3 \\ 0 \end{pmatrix}$  or  $\begin{pmatrix} \vec{u} \\ 0 \end{pmatrix}$



# Homogenous Coordinates

- Let  $M_{3 \times 3}$  be a 3x3 rotation or scaling matrix (as discussed previously)
- Then, the transformation of a point  $\vec{x}$  is given by  $M_{3 \times 3} \vec{x}$

- To produce the same result for  $\begin{pmatrix} \vec{x} \\ 1 \end{pmatrix}$ , use the 4x4 matrix  $\begin{pmatrix} M_{3 \times 3} & \begin{pmatrix} 0 \\ 0 \\ 0 \end{pmatrix} \\ 0 & 0 & 0 & 1 \end{pmatrix} \begin{pmatrix} x \\ y \\ z \\ 1 \end{pmatrix} = \begin{pmatrix} M_{3 \times 3} \vec{x} \\ 1 \end{pmatrix}$

- Similarly, for a vector  $\begin{pmatrix} M_{3 \times 3} & \begin{pmatrix} 0 \\ 0 \\ 0 \end{pmatrix} \\ 0 & 0 & 0 & 1 \end{pmatrix} \begin{pmatrix} u_1 \\ u_2 \\ u_3 \\ 0 \end{pmatrix} = \begin{pmatrix} M_{3 \times 3} \vec{u} \\ 0 \end{pmatrix}$



# Translation

- To translate a point  $\vec{x}$  by some amount  $\vec{t} = \begin{pmatrix} t_1 \\ t_2 \\ t_3 \end{pmatrix}$  one multiplies by a 4x4 matrix


$$\begin{pmatrix} I_{3 \times 3} & \begin{pmatrix} t_1 \\ t_2 \\ t_3 \end{pmatrix} \\ 0 & 0 & 0 & 1 \end{pmatrix} \begin{pmatrix} x \\ y \\ z \\ 1 \end{pmatrix} = \begin{pmatrix} \vec{x} + \vec{t} \\ 1 \end{pmatrix} \text{ where the } 3 \times 3 \text{ identity is } I_{3 \times 3} = \begin{pmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{pmatrix}$$

- For a vector  $\begin{pmatrix} I_{3 \times 3} & \begin{pmatrix} t_1 \\ t_2 \\ t_3 \end{pmatrix} \\ 0 & 0 & 0 & 1 \end{pmatrix} \begin{pmatrix} u_1 \\ u_2 \\ u_3 \\ 0 \end{pmatrix} = \begin{pmatrix} \vec{u} \\ 0 \end{pmatrix}$ , which has no effect (as desired)

- Translations preserves line segments and angles between them, and thus shapes

# Shape is Preserved


- We can translate entire objects without changing them






# Composite Transformations

- Suppose one wants to rotate 45 degrees about the point (1,1)




- These transformations can be multiplied together to get a single matrix  $M = T(1,1)R(45)T(-1,-1)$  that can be used to multiply every relevant point in the entire object:




# Order Matters

- Matrix multiplication does not commute:  $AB \neq BA$
- The rightmost transform is applied to the points first



$T(1,1)R(45)$

$\neq$




$R(45)T(1,1)$



# Hierarchical Transforms

- $M_1$  transforms the teapot from its object space to the table's object space
- $M_2$  transforms the table from its object space to world space
- $M_2M_1$  transforms the teapot from its object space to world space




# Using Transformations

- Create objects (or parts of objects) in convenient coordinate systems
  - Assemble objects from their parts (using transformations)
  - Then, transform the assembled object into the scene (via hierarchical transformations)
  - Can make multiple copies (even of different sizes) of the same object (simply) by adding another transform stack (efficiently, i.e. without creating a new copy of the object)
- 
- Helpful Hint: Always compute composite transforms for objects or sub-objects, and apply the single composite transform to all relevant points (it's a lot faster)
  - Helpful Hint: Orientation is best done first: place the object at the center of the target coordinate system, and rotate it to the desired orientation. Afterwards, translate the object to the correct location.




# Screen Space Projection

- Moving geometry from world space to screen space can create significant distortion
- This is because  $\frac{1}{z}$  is highly nonlinear




$$\frac{x}{z} = \frac{x'}{h}$$



$$x' = h \frac{x}{z}$$

and

$$\frac{y}{z} = \frac{y'}{h}$$



$$y' = h \frac{y}{z}$$




# Matrix Form

- Express the screen space result in homogeneous coordinates as  $\begin{pmatrix} x'w' \\ y'w' \\ z'w' \\ w' \end{pmatrix}$
- Setting  $w' = z$  gives the desired  $\frac{1}{z}$  when dividing by  $w'$
- Consider the following transformation:  $\begin{pmatrix} x'w' \\ y'w' \\ z'w' \\ w' \end{pmatrix} = \begin{pmatrix} h & 0 & 0 & 0 \\ 0 & h & 0 & 0 \\ 0 & 0 & a & b \\ 0 & 0 & 1 & 0 \end{pmatrix} \begin{pmatrix} x \\ y \\ z \\ 1 \end{pmatrix}$
- This has  $w' = z$ ,  $x'w' = hx$  or  $x' = \frac{hx}{z}$ , and  $y'w' = hy$  or  $y' = \frac{hy}{z}$  (as desired)
- Homogenous coordinates allows the nonlinear  $\frac{1}{z}$  to be expressed with linear matrix multiplication (so it can be added to the matrix multiplication stack)!



# Perspective Projection

- The third equation in the linear system is  $z'w' = az + b$  or  $z'z = az + b$ , but  $z$  values are not needed since the projected points all lie on  $z = h$  image plane
- However, computing  $z'$  as a monotonically increasing function of  $z$  allows it to be used to determine occlusions (for alpha channel transparency)
- If  $z = n$  is the near clipping plane and  $z = f$  is the far clipping plane, these clipping planes can be preserved in  $z'$  by setting  $z' = n$  and  $z' = f$  (respectively)
- This gives 2 equations in 2 unknowns:  $n^2 = an + b$  and  $f^2 = af + b$  leads to  $a = n + f$  and  $b = -fn$



- This transforms the viewing frustum into an orthographic volume in screen space